

Assignment 3

ECSE 420: Parallel Computing

Due: December 4, 2018

Submission instructions: Students must submit a pdf file that contains the following information: student's names, student ids, instructions on how to run each file and the associated question solved. Students are also expected to submit a zip file containing their source code. This code must compile and run without error. Code must be well formatted, commented, and follow the google java style guide. For more information, see the ECSE420AssignmentSubmissionInstructions.pdf in the Assignment section on myCourses.

Questions

1. (18 marks) Chapter 7, Memory access, Anderson Lock

The following benchmark is designed to measure the average time for reading array A containing L elements from memory, when the array elements are s words apart:

```
for stride  $s$  from 4 Bytes (1 word) to  $L/2$  by  $2x$ 
  time the following loop (repeat many times and average)
    for  $i$  from 0 to  $L-1$ 
      load  $A[i+s]$  from memory (4 Bytes)
```

The results obtained from the benchmark are shown in Fig.1:




Figure1. Average time for reading each array element

Assuming only one level of cache exists, the cache line size is 4 words and a single processor is running the benchmark, answer the following questions:

1.1 In Fig.1, when the total number of elements of the array - L - is smaller than L' , the average time per access remains constant. What does the value of L' represent? What does time t_0 show?

1.2 When L is larger than L' , what does time t_1 indicate in Fig.1?

1.3 For each part of the graph – parts 1, 2 and 3 -, justify its behaviour.

We know a phenomenon called **false sharing** could cause unnecessary invalidations in ALock (Anderson Lock), and one way to avoid false sharing is to **pad** array elements so that distinct elements are mapped to distinct cache lines.



- 1.4 Considering the results from Fig.1, how could the padding technique used in ALock degrade the overall performance of the lock?
2. (18 marks) Chapter 9, examining the fine-grained algorithm
 - 2.1. Provide the code for the *contains()* method missing from the fine-grained algorithm described in chapter 9.
 - 2.2. Write a test to verify the *contains()* method and explain why your implementation is correct.
3. (12 marks) Chapter 10, designing a bounded lock-based queue
 - 3.1. Design a bounded lock-based queue implementation using an array instead of a linked list. Allow parallelism by using two separate locks for head and tail.
 - 3.2. Try to transform your algorithm to be lock-free. Where do you run into difficulty?
4. (32 marks) Chapter 16, matrix vector multiplication
 - 4.1. Implement a sequential matrix vector multiplication.
 - 4.2. Give an efficient and highly parallel multithreaded algorithm for multiplying an $n \times n$ matrix A by a length- n vector x that achieves work $\Theta(n^2)$ and critical path $\Theta(\log n)$.
 - 4.3. Write a test program that measures the execution time for multiplying a 2,000 by 2,000 matrix with a corresponding 2000 wide vector using the parallel method and sequential method, respectively. Discuss the execution time of the two methods and compute the speedup. Be sure to indicate how many threads are being used.
 - 4.4. Analyze and discuss the work and critical-path length of your implementation, and give the *parallelism*.

Total: 80 marks

